

ASSGN_3_PART_1

February 20, 2022

```
[1]: # MLFA Assignment 3
```

```
# Team Members -  
# Deepansh Agrawal - 19MI10018  
# Rohit Ranjan - 20CS30066  
# Neha Gupta - 20CH10094  
# Gautam Jaju - 20AG30015  
# Siddharth Madhupati - 20ME30083  
# Madiha Hanifa - 20MF10018  
# Himadri Pandya - 20ME10047
```

```
[2]: # import commands  
import numpy as np  
import matplotlib.pyplot as plt  
import random  
import tensorflow as tf  
from tensorflow import keras  
from keras.layers.core import Dense, Dropout, Activation  
from keras.layers import BatchNormalization  
from keras.utils import np_utils  
  
# loading the dataset  
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```

```
[3]: # reshaping the input data to the range 0 - 1  
x_train = x_train.reshape(-1, 28 * 28).astype("float32") / 255.0  
x_test = x_test.reshape(-1, 28 * 28).astype("float32") / 255.0
```

```
[4]: print(x_train.shape)  
print(y_train.shape)
```

```
(60000, 784)  
(60000,)
```

```
[5]: # Implementation of the uniform network  
# There are 88 neurons in each layer except the first layer and the output layer  
# The parameters of the input layer (39250) and first layer (4488) are  
# ↪ discounted
```

```
# The number of parameters is constant after that
# The number of parameters is different in the last hidden layer as specified,
↳ in the problem
```

```
[6]: initializer = tf.keras.initializers.GlorotNormal()
UniformNet = keras.Sequential([
    keras.layers.InputLayer((784)),
    keras.layers.Dense(50, activation='relu',
↳ name='first_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='second_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='third_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='fourth_layer', kernel_initializer=initializer), Dropout(0.3),
↳ BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='fifth_layer', kernel_initializer=initializer), Dropout(0.3),
↳ BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='sixth_layer', kernel_initializer=initializer), Dropout(0.3),
↳ BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='seventh_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='eighth_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(88, activation='relu',
↳ name='nineth_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(10, activation='softmax',
↳ name='output_layer', kernel_initializer=initializer)
])
```

```
[7]: UniformNet.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
first_layer (Dense)	(None, 50)	39250
batch_normalization (Batch Normalization)	(None, 50)	200
second_layer (Dense)	(None, 88)	4488
batch_normalization_1 (Batch Normalization)	(None, 88)	352

third_layer (Dense)	(None, 88)	7832
batch_normalization_2 (Batch Normalization)	(None, 88)	352
fourth_layer (Dense)	(None, 88)	7832
dropout (Dropout)	(None, 88)	0
batch_normalization_3 (Batch Normalization)	(None, 88)	352
fifth_layer (Dense)	(None, 88)	7832
dropout_1 (Dropout)	(None, 88)	0
batch_normalization_4 (Batch Normalization)	(None, 88)	352
sixth_layer (Dense)	(None, 88)	7832
dropout_2 (Dropout)	(None, 88)	0
batch_normalization_5 (Batch Normalization)	(None, 88)	352
seventh_layer (Dense)	(None, 88)	7832
batch_normalization_6 (Batch Normalization)	(None, 88)	352
eighth_layer (Dense)	(None, 88)	7832
batch_normalization_7 (Batch Normalization)	(None, 88)	352
ninth_layer (Dense)	(None, 88)	7832
batch_normalization_8 (Batch Normalization)	(None, 88)	352
output_layer (Dense)	(None, 10)	890

```
=====
Total params: 102,468
Trainable params: 100,960
Non-trainable params: 1,508
```

```
-----  
[8]: UniformNet.compile(  
    loss=keras.losses.SparseCategoricalCrossentropy(),  
    optimizer=keras.optimizers.Adam(learning_rate=0.0001),  
    metrics=["accuracy"],  
)
```

```
[9]: UniformNet.fit(x_train, y_train, batch_size=128, epochs=20, verbose=1)  
UniformNet.evaluate(x_test, y_test, batch_size=128, verbose=1)
```

```
Epoch 1/20  
469/469 [=====] - 10s 15ms/step - loss: 2.3375 -  
accuracy: 0.2103  
Epoch 2/20  
469/469 [=====] - 4s 9ms/step - loss: 1.6245 -  
accuracy: 0.4308  
Epoch 3/20  
469/469 [=====] - 4s 9ms/step - loss: 1.1427 -  
accuracy: 0.6094  
Epoch 4/20  
469/469 [=====] - 4s 9ms/step - loss: 0.8326 -  
accuracy: 0.7289  
Epoch 5/20  
469/469 [=====] - 4s 8ms/step - loss: 0.6374 -  
accuracy: 0.8036  
Epoch 6/20  
469/469 [=====] - 4s 9ms/step - loss: 0.5170 -  
accuracy: 0.8464  
Epoch 7/20  
469/469 [=====] - 4s 9ms/step - loss: 0.4356 -  
accuracy: 0.8745  
Epoch 8/20  
469/469 [=====] - 4s 9ms/step - loss: 0.3830 -  
accuracy: 0.8918  
Epoch 9/20  
469/469 [=====] - 4s 9ms/step - loss: 0.3396 -  
accuracy: 0.9048  
Epoch 10/20  
469/469 [=====] - 4s 9ms/step - loss: 0.3092 -  
accuracy: 0.9142  
Epoch 11/20  
469/469 [=====] - 4s 9ms/step - loss: 0.2814 -  
accuracy: 0.9223  
Epoch 12/20  
469/469 [=====] - 4s 8ms/step - loss: 0.2645 -  
accuracy: 0.9281  
Epoch 13/20
```

```

469/469 [=====] - 4s 8ms/step - loss: 0.2478 -
accuracy: 0.9330
Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.2300 -
accuracy: 0.9379
Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.2185 -
accuracy: 0.9418
Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.2094 -
accuracy: 0.9442
Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.1960 -
accuracy: 0.9481
Epoch 18/20
469/469 [=====] - 4s 8ms/step - loss: 0.1858 -
accuracy: 0.9499
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.1771 -
accuracy: 0.9525
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.1685 -
accuracy: 0.9547
79/79 [=====] - 1s 3ms/step - loss: 0.1397 - accuracy:
0.9632

```

```
[9]: [0.1396992951631546, 0.9631999731063843]
```

```

[10]: # Implementation of the pyramid network
      # There are varying number of neurons in each layer
      # The parameters of the input layer (39250) and first layer (200) are discounted
      # The number of parameters falls by a factor of two after that
      # The number of parameters is different in the last hidden layer as specified
      → in the problem

```

```

[11]: initializer = tf.keras.initializers.GlorotNormal()
      PyramidNet = keras.Sequential([
          keras.layers.InputLayer((784)),
          keras.layers.Dense(50, activation='relu',
      → name='first_layer', kernel_initializer=initializer), BatchNormalization(),
          keras.layers.Dense(16, activation='relu',
      → name='second_layer', kernel_initializer=initializer), BatchNormalization(),
          keras.layers.Dense(1250, activation='relu',
      → name='third_layer', kernel_initializer=initializer), BatchNormalization(),
          keras.layers.Dense(8, activation='relu',
      → name='fourth_layer', kernel_initializer=initializer), Dropout(0.3),
      → BatchNormalization(),

```

```

        keras.layers.Dense(625, activation='relu',
↪name='fifth_layer', kernel_initializer=initializer), Dropout(0.3),
↪BatchNormalization(),
        keras.layers.Dense(4, activation='relu',
↪name='sixth_layer', kernel_initializer=initializer), Dropout(0.3),
↪BatchNormalization(),
        keras.layers.Dense(312, activation='relu',
↪name='seventh_layer', kernel_initializer=initializer), BatchNormalization(),
        keras.layers.Dense(2, activation='relu',
↪name='eighth_layer', kernel_initializer=initializer), BatchNormalization(),
        keras.layers.Dense(156, activation='relu',
↪name='nineth_layer', kernel_initializer=initializer), BatchNormalization(),
        keras.layers.Dense(10, activation='softmax',
↪name='output_layer', kernel_initializer=initializer)
])

```

[12]: `PyramidNet.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
first_layer (Dense)	(None, 50)	39250
batch_normalization_9 (Batch Normalization)	(None, 50)	200
second_layer (Dense)	(None, 16)	816
batch_normalization_10 (Batch Normalization)	(None, 16)	64
third_layer (Dense)	(None, 1250)	21250
batch_normalization_11 (Batch Normalization)	(None, 1250)	5000
fourth_layer (Dense)	(None, 8)	10008
dropout_3 (Dropout)	(None, 8)	0
batch_normalization_12 (Batch Normalization)	(None, 8)	32
fifth_layer (Dense)	(None, 625)	5625
dropout_4 (Dropout)	(None, 625)	0

batch_normalization_13 (Batch Normalization)	(None, 625)	2500
sixth_layer (Dense)	(None, 4)	2504
dropout_5 (Dropout)	(None, 4)	0
batch_normalization_14 (Batch Normalization)	(None, 4)	16
seventh_layer (Dense)	(None, 312)	1560
batch_normalization_15 (Batch Normalization)	(None, 312)	1248
eighth_layer (Dense)	(None, 2)	626
batch_normalization_16 (Batch Normalization)	(None, 2)	8
nineth_layer (Dense)	(None, 156)	468
batch_normalization_17 (Batch Normalization)	(None, 156)	624
output_layer (Dense)	(None, 10)	1570

```
=====
Total params: 93,369
Trainable params: 88,523
Non-trainable params: 4,846
-----
```

```
[13]: PyramidNet.compile(
        loss=keras.losses.SparseCategoricalCrossentropy(),
        optimizer=keras.optimizers.Adam(learning_rate=0.0001),
        metrics=["accuracy"],
    )
```

```
[14]: PyramidNet.fit(x_train, y_train, batch_size=128, epochs=20, verbose=1)
PyramidNet.evaluate(x_test, y_test, batch_size=128, verbose=1)
```

```
Epoch 1/20
469/469 [=====] - 10s 16ms/step - loss: 2.2005 -
accuracy: 0.1786
Epoch 2/20
469/469 [=====] - 7s 16ms/step - loss: 1.8831 -
```

```

accuracy: 0.2492
Epoch 3/20
469/469 [=====] - 7s 15ms/step - loss: 1.6977 -
accuracy: 0.3248
Epoch 4/20
469/469 [=====] - 7s 16ms/step - loss: 1.5016 -
accuracy: 0.3735
Epoch 5/20
469/469 [=====] - 7s 15ms/step - loss: 1.4055 -
accuracy: 0.4181
Epoch 6/20
469/469 [=====] - 7s 15ms/step - loss: 1.3172 -
accuracy: 0.4679
Epoch 7/20
469/469 [=====] - 7s 15ms/step - loss: 1.2528 -
accuracy: 0.5031
Epoch 8/20
469/469 [=====] - 7s 15ms/step - loss: 1.2146 -
accuracy: 0.5209
Epoch 9/20
469/469 [=====] - 7s 15ms/step - loss: 1.1708 -
accuracy: 0.5349
Epoch 10/20
469/469 [=====] - 7s 16ms/step - loss: 1.1302 -
accuracy: 0.5505
Epoch 11/20
469/469 [=====] - 7s 15ms/step - loss: 1.0979 -
accuracy: 0.5633
Epoch 12/20
469/469 [=====] - 7s 15ms/step - loss: 1.0546 -
accuracy: 0.5766
Epoch 13/20
469/469 [=====] - 7s 15ms/step - loss: 1.0418 -
accuracy: 0.5830
Epoch 14/20
469/469 [=====] - 7s 15ms/step - loss: 1.0223 -
accuracy: 0.5849
Epoch 15/20
469/469 [=====] - 7s 15ms/step - loss: 0.9996 -
accuracy: 0.5920
Epoch 16/20
469/469 [=====] - 7s 15ms/step - loss: 0.9871 -
accuracy: 0.5903
Epoch 17/20
469/469 [=====] - 7s 16ms/step - loss: 0.9787 -
accuracy: 0.5960
Epoch 18/20
469/469 [=====] - 7s 16ms/step - loss: 0.9618 -

```



```

accuracy: 0.5973
Epoch 19/20
469/469 [=====] - 7s 16ms/step - loss: 0.9544 -
accuracy: 0.6003
Epoch 20/20
469/469 [=====] - 7s 16ms/step - loss: 0.9478 -
accuracy: 0.5988
79/79 [=====] - 1s 4ms/step - loss: 2.3110 - accuracy:
0.2793

```

```
[14]: [2.311028480529785, 0.2793000042438507]
```

```
[15]: # Implementation of the pyramid network
# There are varying number of neurons in each layer
# The parameters of the input layer (39250) and first layer (255) are discounted
# The number of parameters increases by a factor of two after that
# The number of parameters is different in the last hidden layer as specified
    ↪ in the problem
```

```
[16]: initializer = tf.keras.initializers.GlorotNormal()
InvPyramidNet = keras.Sequential([
    keras.layers.InputLayer((784)),
    keras.layers.Dense(50, activation='relu',
    ↪ name='first_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(5, activation='relu',
    ↪ name='second_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(94, activation='relu',
    ↪ name='third_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(10, activation='relu',
    ↪ name='fourth_layer', kernel_initializer=initializer), Dropout(0.3),
    ↪ BatchNormalization(),
    keras.layers.Dense(188, activation='relu',
    ↪ name='fifth_layer', kernel_initializer=initializer), Dropout(0.3),
    ↪ BatchNormalization(),
    keras.layers.Dense(20, activation='relu',
    ↪ name='sixth_layer', kernel_initializer=initializer), Dropout(0.3),
    ↪ BatchNormalization(),
    keras.layers.Dense(376, activation='relu',
    ↪ name='seventh_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(40, activation='relu',
    ↪ name='eighth_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(750, activation='relu',
    ↪ name='ninth_layer', kernel_initializer=initializer), BatchNormalization(),
    keras.layers.Dense(10, activation='softmax',
    ↪ name='output_layer', kernel_initializer=initializer)
])
```

```
[17]: InvPyramidNet.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
first_layer (Dense)	(None, 50)	39250
batch_normalization_18 (Batch Normalization)	(None, 50)	200
second_layer (Dense)	(None, 5)	255
batch_normalization_19 (Batch Normalization)	(None, 5)	20
third_layer (Dense)	(None, 94)	564
batch_normalization_20 (Batch Normalization)	(None, 94)	376
fourth_layer (Dense)	(None, 10)	950
dropout_6 (Dropout)	(None, 10)	0
batch_normalization_21 (Batch Normalization)	(None, 10)	40
fifth_layer (Dense)	(None, 188)	2068
dropout_7 (Dropout)	(None, 188)	0
batch_normalization_22 (Batch Normalization)	(None, 188)	752
sixth_layer (Dense)	(None, 20)	3780
dropout_8 (Dropout)	(None, 20)	0
batch_normalization_23 (Batch Normalization)	(None, 20)	80
seventh_layer (Dense)	(None, 376)	7896
batch_normalization_24 (Batch Normalization)	(None, 376)	1504

eighth_layer (Dense)	(None, 40)	15080
batch_normalization_25 (Batch Normalization)	(None, 40)	160
ninth_layer (Dense)	(None, 750)	30750
batch_normalization_26 (Batch Normalization)	(None, 750)	3000
output_layer (Dense)	(None, 10)	7510

```

=====
Total params: 114,235
Trainable params: 111,169
Non-trainable params: 3,066
-----

```

```
[18]: InvPyramidNet.compile(
        loss=keras.losses.SparseCategoricalCrossentropy(),
        optimizer=keras.optimizers.Adam(learning_rate=0.0001),
        metrics=["accuracy"],
    )
```

```
[19]: InvPyramidNet.fit(x_train, y_train, batch_size=128, epochs=20, verbose=1)
      InvPyramidNet.evaluate(x_test, y_test, batch_size=128, verbose=1)
```

```

Epoch 1/20
469/469 [=====] - 8s 13ms/step - loss: 2.1832 -
accuracy: 0.2240
Epoch 2/20
469/469 [=====] - 6s 12ms/step - loss: 1.4959 -
accuracy: 0.4511
Epoch 3/20
469/469 [=====] - 6s 12ms/step - loss: 1.1774 -
accuracy: 0.5608
Epoch 4/20
469/469 [=====] - 6s 12ms/step - loss: 1.0142 -
accuracy: 0.6153
Epoch 5/20
469/469 [=====] - 6s 12ms/step - loss: 0.9120 -
accuracy: 0.6561
Epoch 6/20
469/469 [=====] - 6s 13ms/step - loss: 0.8488 -
accuracy: 0.6820
Epoch 7/20
469/469 [=====] - 6s 13ms/step - loss: 0.7879 -
accuracy: 0.6998

```

```

Epoch 8/20
469/469 [=====] - 6s 13ms/step - loss: 0.7367 -
accuracy: 0.7251
Epoch 9/20
469/469 [=====] - 6s 13ms/step - loss: 0.7038 -
accuracy: 0.7403
Epoch 10/20
469/469 [=====] - 6s 12ms/step - loss: 0.6601 -
accuracy: 0.7586
Epoch 11/20
469/469 [=====] - 6s 12ms/step - loss: 0.6385 -
accuracy: 0.7674
Epoch 12/20
469/469 [=====] - 6s 13ms/step - loss: 0.5991 -
accuracy: 0.7829
Epoch 13/20
469/469 [=====] - 6s 13ms/step - loss: 0.5774 -
accuracy: 0.7957
Epoch 14/20
469/469 [=====] - 6s 12ms/step - loss: 0.5558 -
accuracy: 0.8032
Epoch 15/20
469/469 [=====] - 6s 12ms/step - loss: 0.5327 -
accuracy: 0.8132
Epoch 16/20
469/469 [=====] - 6s 12ms/step - loss: 0.5107 -
accuracy: 0.8351
Epoch 17/20
469/469 [=====] - 6s 12ms/step - loss: 0.4901 -
accuracy: 0.8507
Epoch 18/20
469/469 [=====] - 6s 13ms/step - loss: 0.4682 -
accuracy: 0.8612
Epoch 19/20
469/469 [=====] - 6s 12ms/step - loss: 0.4508 -
accuracy: 0.8694
Epoch 20/20
469/469 [=====] - 6s 12ms/step - loss: 0.4383 -
accuracy: 0.8735
79/79 [=====] - 1s 3ms/step - loss: 1.0184 - accuracy:
0.6292

```

[19]: [1.0183591842651367, 0.6291999816894531]

[19]: