

# Assignment\_5

March 22, 2022

```
[ ]: # Rohit Ranjan  
# 20CS30066  
# Assignment 5
```

```
[1]: #import commands  
import numpy as np  
import matplotlib.pyplot as plt  
plt.rcParams["figure.figsize"] = (80,80)  
import pandas as pd  
from sklearn.metrics import mean_squared_error  
from sklearn.model_selection import GridSearchCV  
  
#xgboost  
import xgboost as xgb  
  
#SVM  
from sklearn.svm import SVR
```

```
[2]: #datafile provided  
filename = "S_n_I_A_N_P_An_Io_noaa2.txt"
```

```
[3]: #preprocessing the datafile  
rows=[]  
for line in open(filename):  
    curr_row = list(map(float, line.split()))  
    rows.append(curr_row)
```

```
[4]: len(rows)
```

```
[4]: 870
```

```
[5]: df = pd.DataFrame(rows)  
df.columns = ['SAM', 'nino', 'ISMR', 'AMO', 'NAO', 'PDO', 'At-nino', 'IOD']
```

```
[6]: cols_ordered = ['SAM', 'nino', 'AMO', 'NAO', 'PDO', 'At-nino', 'IOD', 'ISMR']  
df = df[cols_ordered]
```

```
[7]: df.head(5)
```

```
[7]:      SAM      nino      AMO      NAO      PDO      At-nino      IOD  \
0  0.599257  0.167651  0.484010  0.000837  1.084100  0.560382 -0.295856
1  0.663879  0.212788  0.603261 -0.317969 -1.465790  1.189540 -0.032952
2 -0.428772  0.336627  0.609364  0.423296  0.114019  0.783738 -0.644708
3  0.513283  0.575112  0.219494  0.352545  1.807430  0.560395 -0.295477
4 -0.260801  0.647786  0.040233 -0.528239 -0.108867  0.727379  0.044719

      ISMR
0  10.24380
1  44.16110
2   4.76807
3 -63.58770
4  12.57720
```

```
[8]: #Experiments with xgBoost
```

```
[9]: true_values = []
pred_values = []
rmse_values = []

for lag in range(1,11):
    # forming features using lag
    y = df['ISMR'][lag:]
    X = df.iloc[:-1*lag,:]

    # converting to efficient data structure
    data_dmatrix = xgb.DMatrix(data=X,label=y)

    # intialising xgboost regressor object and fitting on data
    xg_reg = xgb.XGBRegressor(colsample_bytree = 0.3, learning_rate = 0.
→1,max_depth = 5, alpha = 10, n_estimators = 1000)
    xg_reg.fit(X,y)

    #generating predictions
    preds = xg_reg.predict(X)

    # calculating error to help as compare lag performance
    rmse = np.sqrt(mean_squared_error(y, preds))

    # appending values to a list for later use
    true_values.append(y)
    pred_values.append(preds)
    rmse_values.append((rmse))
```

```
[10]: rmse_table = pd.DataFrame(list(np.array(rmse_values).
    ↳ reshape((len(rmse_values),-1)).T))
rmse_table.columns = ["Lag_"+str(i)) for i in range(1,11)]
```

```
[11]: rmse_table
```

```
[11]:      Lag_1      Lag_2      Lag_3      Lag_4      Lag_5      Lag_6      Lag_7 \
0  1.586117  1.700843  1.797371  1.75844  1.742344  1.723301  1.743826

      Lag_8      Lag_9      Lag_10
0  1.765855  1.814621  1.634117
```

```
[12]: # computing table for different lags
table = pd.DataFrame()
table['True ISMR'] = df['ISMR']
```











```
[13]: # padding the lag columns
for i in range(len(pred_values)):
    pred_lag = list(pred_values[i])
    for i in range(df.shape[0]-len(pred_lag)):
        pred_lag.insert(0,np.NaN)
    table['Lag_'+str(i+1)] = pred_lag
```

```
[14]: table[0:20]
```

```
[14]:      True ISMR      Lag_1      Lag_2      Lag_3      Lag_4      Lag_5 \
0    10.243800      NaN      NaN      NaN      NaN      NaN
1    44.161100  43.150269      NaN      NaN      NaN      NaN
2     4.768070   5.338588   4.439160      NaN      NaN      NaN
3   -63.587700 -62.148609 -62.418106 -58.706940      NaN      NaN
4    12.577200  12.076989  11.821412  11.542588  13.834686      NaN
5   -41.577200 -41.952026 -38.927246 -38.944092 -42.661339 -38.924084
6    -9.546880  -9.879082 -10.141637  -9.084824  -9.147939  -9.122007
7    25.370500  24.726170  23.547773  23.525265  24.865967  24.532143
8    18.277400  18.698280  17.388618  16.878601  20.062557  17.476351
9     2.221670   1.521839   2.148856   3.221498   1.718592   0.318806
10   16.886600  17.196455  16.851954  17.330164  16.662622  16.647011
11    0.132196   0.001765   1.290223   2.877347   0.691788   0.578885
12  -10.537500  -9.827923  -9.441161 -12.002758  -9.992893 -10.177773
13  -50.820200 -48.424641 -46.539761 -48.758595 -48.716595 -50.296795
14   -8.513200  -6.482193  -8.650455  -8.616445  -8.550304  -7.871543
15  -28.769000 -27.899918 -25.707705 -28.510756 -27.202774 -27.854412
16   -5.404050  -5.391354  -4.667356  -5.478644  -4.768011  -5.960122
17  -17.658400 -17.083843 -17.908487 -16.732082 -15.841544 -17.970730
18   14.971800  16.190086  14.987142  14.808299  14.542233  14.397070
19   64.089200  62.903187  64.037880  62.913769  63.030930  62.196003
```

|    | Lag_6      | Lag_7      | Lag_8      | Lag_9      | Lag_10     |
|----|------------|------------|------------|------------|------------|
| 0  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 1  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 2  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 3  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 4  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 5  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 6  | -7.095208  | NaN        | NaN        | NaN        | NaN        |
| 7  | 24.191381  | 25.250797  | NaN        | NaN        | NaN        |
| 8  | 16.491549  | 18.776773  | 17.648241  | NaN        | NaN        |
| 9  | 1.528083   | 3.899342   | 1.213346   | 1.961416   | NaN        |
| 10 | 17.723703  | 16.984150  | 16.306791  | 16.890469  | 17.297546  |
| 11 | 1.266631   | 1.552985   | 2.110994   | 1.315464   | 0.438608   |
| 12 | -11.356400 | -11.348008 | -10.918290 | -11.448859 | -10.131416 |
| 13 | -47.985355 | -46.207077 | -48.974602 | -47.702118 | -49.232628 |
| 14 | -7.291180  | -7.087028  | -7.215755  | -8.229162  | -8.708284  |
| 15 | -26.521212 | -25.365711 | -28.066278 | -26.121048 | -26.770824 |
| 16 | -6.601791  | -5.548816  | -4.540854  | -5.208636  | -4.742649  |
| 17 | -17.766972 | -15.611817 | -16.529551 | -17.022909 | -16.926622 |
| 18 | 14.450985  | 14.158209  | 13.296273  | 14.539018  | 13.520664  |
| 19 | 63.259460  | 64.154404  | 62.819069  | 61.924686  | 59.933323  |

```
[15]: # plotting scatter plots
fig, axs = plt.subplots(10)
for i in range(10):
    axs[i].scatter(true_values[i], pred_values[i])
    axs[i].xlabel = 'True Values'
    axs[i].ylabel = 'Predicted Values'
    axs[i].set_title("Lag = "+str(i+1))
    axs[i].set_aspect('equal', adjustable='box')
```



```
[16]: # The best lag in drivers is 1 from our analysis. From the plotted graphs that
      ↳ xgboost trained with lag=1 dataset gives preds that most closely resemble
      ↳ the true values. This leads to the graph being closest to a straight line
      ↳ among all 10 lags.
```

```
[17]: # SVM - Linear Kernel
```

```
[18]: true_values = []
      pred_values = []
      rmse_values = []
      for lag in range(1,11):
          # forming features using lag
          y = df['ISMR'][lag:]
          X = df.iloc[:-1*lag,:]

          # initialising SVM regressor object and fitting on data
          svr = SVR(kernel='linear', C=1.0, epsilon=0.1)
          svr.fit(X,y)

          #generating predictions
          preds = svr.predict(X)

          # calculating error to help as compare lag performance
          rmse = np.sqrt(mean_squared_error(y, preds))

          # appending values to a list for later use
          true_values.append(y)
          pred_values.append(preds)
          rmse_values.append((rmse))
```

```
[19]: rmse_table = pd.DataFrame(list(np.array(rmse_values).
      ↳ reshape((len(rmse_values),-1)).T))
      rmse_table.columns = [("Lag_"+str(i)) for i in range(1,11)]
```

```
[20]: rmse_table
```

```
[20]:      Lag_1      Lag_2      Lag_3      Lag_4      Lag_5      Lag_6      Lag_7 \
0  31.550819  32.17455  32.850646  33.040213  32.987998  32.984692  32.878436

      Lag_8      Lag_9      Lag_10
0  32.752821  32.840307  32.994115
```

```
[21]: table = pd.DataFrame()
      table['True ISMR'] = df['ISMR']
```

```
[22]: for i in range(len(pred_values)):
      pred_lag = list(pred_values[i])
      for i in range(df.shape[0]-len(pred_lag)):
          pred_lag.insert(0,np.NaN)
          table['Lag_'+str(i+1)] = pred_lag
```

```
[23]: table[0:20]
```

```
[23]:   True ISMR      Lag_1      Lag_2      Lag_3      Lag_4      Lag_5      Lag_6  \
0   10.243800         NaN         NaN         NaN         NaN         NaN         NaN
1   44.161100    0.823180         NaN         NaN         NaN         NaN         NaN
2    4.768070    4.867730    1.860854         NaN         NaN         NaN         NaN
3  -63.587700   -2.126838    6.851839    1.727825         NaN         NaN         NaN
4   12.577200  -11.990512    0.269369    3.652074    5.455071         NaN         NaN
5  -41.577200   -5.372353   -8.518099    1.674918    5.052237    2.126606         NaN
6   -9.546880   -5.253140   -3.546155   -1.582393    4.772103    6.275801    3.273256
7   25.370500    6.234563   -5.974274   -1.381092   -0.262685    4.042350    2.506127
8   18.277400    8.497448    2.104331   -3.466746    1.935317    0.441458    7.155424
9    2.221670   -3.681849    7.646596   -0.605787   -8.325702    1.718356    5.751010
10  16.886600    6.385370   -3.501795    0.865078   -0.788679   -0.207964    2.895361
11    0.132196    5.670740    7.593964   -0.339189   -1.754547   -2.423448    1.748274
12 -10.537500    9.512690    7.718172    5.689226    5.991894    0.598983   -0.125172
13 -50.820200    8.927082    6.986084    3.283073    4.447418    2.450394   -3.112934
14   -8.513200    1.763999    5.109834   -1.408233    1.035883   -2.301114    6.074265
15 -28.769000   -5.630880   -1.378075    2.045129   -6.512921   -0.242641    0.451446
16   -5.404050   -9.985943   -3.864202   -2.316010   -0.711510   -3.434859   -4.223349
17 -17.658400   -9.433855   -7.045584   -1.799386   -8.330568   -3.010712   -7.293407
18  14.971800    2.676125   -1.344234   -2.299299    0.999803    1.494560   -4.691778
19  64.089200   12.543811    2.555395    2.245494    4.976250   -0.141671   -2.489480

      Lag_7      Lag_8      Lag_9      Lag_10
0         NaN         NaN         NaN         NaN
1         NaN         NaN         NaN         NaN
2         NaN         NaN         NaN         NaN
3         NaN         NaN         NaN         NaN
4         NaN         NaN         NaN         NaN
5         NaN         NaN         NaN         NaN
6         NaN         NaN         NaN         NaN
7    1.702393         NaN         NaN         NaN
8    3.645954    2.205893         NaN         NaN
9    4.408004   -0.365594    0.525703         NaN
10   3.330357    3.415228   -0.883780    4.134303
11   0.031863    2.956235    1.789533   -1.509402
12  -0.076477    1.646860    0.458877    2.655329
13  -3.220253   -1.874773    0.120246    6.220969
14  -0.375654   -1.772087   -4.166345   -1.135234
15   0.307970   -4.446781   -4.877276   -7.316638
```

```


16 -2.660224  4.907276 -3.198715 -1.138683
17 -1.765303 -1.450087  1.270269 -4.213599
18 -4.305335 -5.159260 -4.342382  1.460853
19 -2.825177 -9.688401 -5.017985  1.351116

```

```

[24]: fig, axs = plt.subplots(10)
      for i in range(10):
          axs[i].scatter(true_values[i],pred_values[i])
          axs[i].xlabel = 'True Values'
          axs[i].ylabel = 'Predicted Values'
          axs[i].set_title("Lag = "+str(i+1))
          axs[i].set_aspect('equal', adjustable='box')

```





```
[25]: # SVM - Poly Kernel
```

```
[26]: true_values = []
pred_values = []
rmse_values = []
for lag in range(1,11):
    # forming features using lag
    y = df['ISMR'][lag:]
    X = df.iloc[:-1*lag,:].

    # intialising SVM regressor object and fitting on data
    svr = SVR(kernel='poly', C=1.0, epsilon=0.1)
    svr.fit(X,y)

    #generating predictions
    preds = svr.predict(X)

    # calculating error to help as compare lag performance
    rmse = np.sqrt(mean_squared_error(y, preds))

    # appending values to a list for later use
    true_values.append(y)
    pred_values.append(preds)
    rmse_values.append((rmse))
```

```
[27]: rmse_table = pd.DataFrame(list(np.array(rmse_values).
    ↪reshape((len(rmse_values),-1)).T))
rmse_table.columns = [f'Lag_{i}' for i in range(1,11)]
```

```
[28]: rmse_table
```

```
[28]:
```

|   | Lag_1     | Lag_2     | Lag_3     | Lag_4     | Lag_5     | Lag_6     | Lag_7     | \ |
|---|-----------|-----------|-----------|-----------|-----------|-----------|-----------|---|
| 0 | 32.611828 | 32.92766  | 33.057109 | 33.041273 | 33.080376 | 32.981059 | 32.938083 |   |
|   |           |           |           |           |           |           |           |   |
|   | Lag_8     | Lag_9     | Lag_10    |           |           |           |           |   |
| 0 | 33.0812   | 33.088109 | 33.141695 |           |           |           |           |   |

```
[29]: table = pd.DataFrame()
table['True ISMR'] = df['ISMR']
```

```
[30]: for i in range(len(pred_values)):
    pred_lag = list(pred_values[i])
    for i in range(df.shape[0]-len(pred_lag)):
        pred_lag.insert(0,np.NaN)
```

```
table['Lag_'+str(i+1)] = pred_lag
```

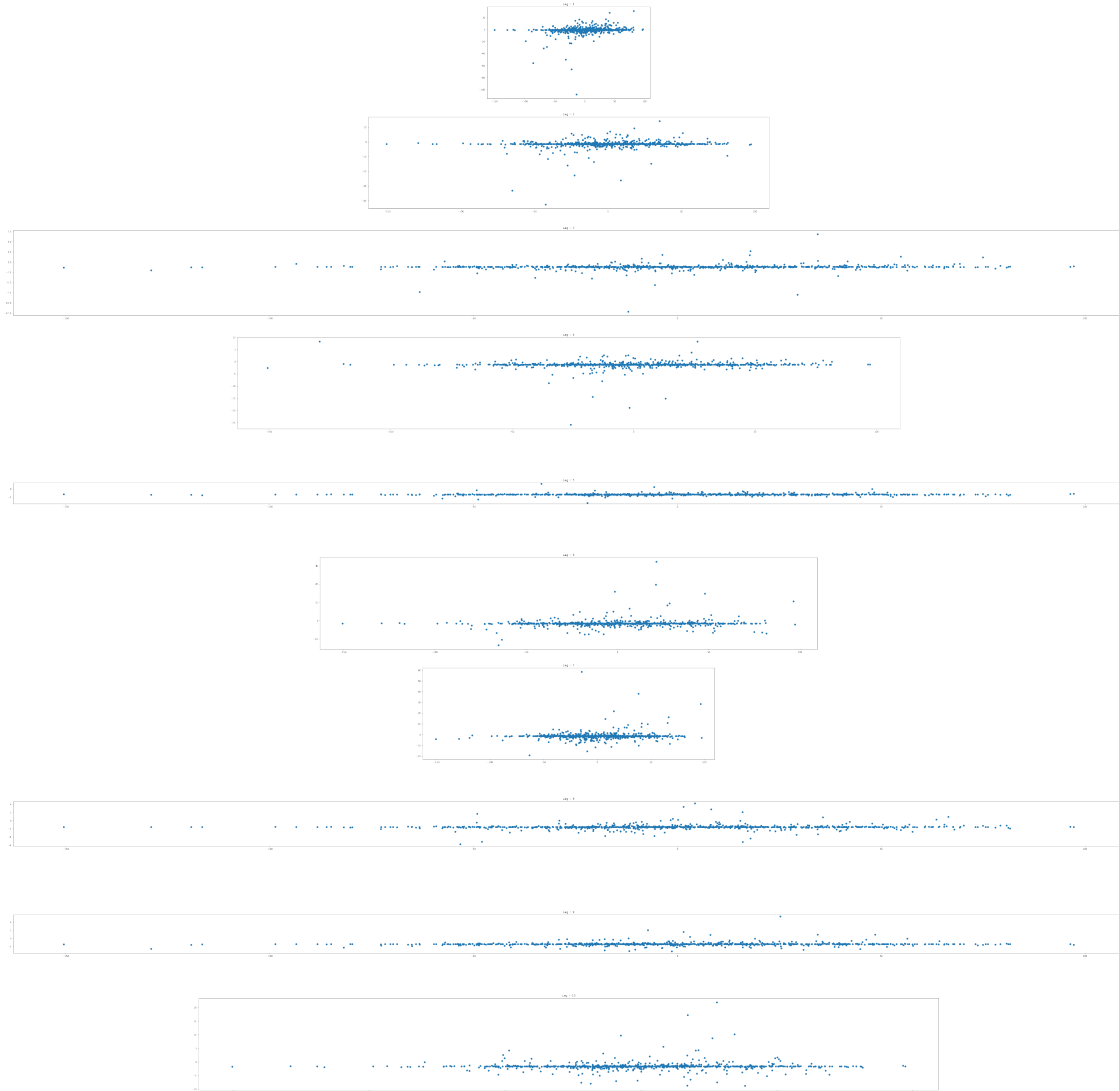
```
[31]: table[0:20]
```

```
[31]:
```

|    | True ISMR  | Lag_1     | Lag_2     | Lag_3     | Lag_4     | Lag_5     | Lag_6     | \ |
|----|------------|-----------|-----------|-----------|-----------|-----------|-----------|---|
| 0  | 10.243800  | NaN       | NaN       | NaN       | NaN       | NaN       | NaN       |   |
| 1  | 44.161100  | -0.553799 | NaN       | NaN       | NaN       | NaN       | NaN       |   |
| 2  | 4.768070   | 1.409998  | -1.337070 | NaN       | NaN       | NaN       | NaN       |   |
| 3  | -63.587700 | -0.594516 | -0.480561 | -1.145828 | NaN       | NaN       | NaN       |   |
| 4  | 12.577200  | -7.831394 | -1.345313 | -1.116867 | -1.147418 | NaN       | NaN       |   |
| 5  | -41.577200 | -0.567826 | -4.416571 | -1.159616 | -0.818578 | -1.357415 | NaN       |   |
| 6  | -9.546880  | -3.268182 | -1.360824 | -1.322785 | -1.161940 | -1.061223 | -1.489415 |   |
| 7  | 25.370500  | -0.600268 | -2.095118 | -1.140791 | -2.427974 | -1.350311 | -1.729758 |   |
| 8  | 18.277400  | -0.185252 | -1.334301 | -1.457471 | -1.146612 | -1.896288 | -1.498301 |   |
| 9  | 2.221670   | -0.424751 | -1.026527 | -1.151829 | -1.736113 | -1.343231 | 1.875892  |   |
| 10 | 16.886600  | -0.594682 | -1.342886 | -1.237701 | -1.156892 | -1.219041 | -1.478619 |   |
| 11 | 0.132196   | -0.451335 | -1.348417 | -1.052843 | -1.133828 | -1.365148 | -0.743560 |   |
| 12 | -10.537500 | -0.596784 | -1.318907 | -1.160308 | -1.054258 | -1.285997 | -1.492600 |   |
| 13 | -50.820200 | -0.601772 | -1.348542 | -1.159339 | -1.164988 | -1.365201 | -1.695292 |   |
| 14 | -8.513200  | -5.034933 | -1.370768 | -1.161943 | -1.177999 | -1.351717 | -1.459198 |   |
| 15 | -28.769000 | -0.615795 | -2.502595 | -1.148397 | -1.166002 | -1.356384 | -1.500031 |   |
| 16 | -5.404050  | -1.097010 | -1.359551 | -2.252933 | -1.173027 | -1.350597 | -1.476281 |   |
| 17 | -17.658400 | -0.606590 | -1.620088 | -1.155168 | -2.731447 | -1.363818 | -1.500766 |   |
| 18 | 14.971800  | -0.776900 | -1.358475 | -1.093639 | -1.165094 | -1.171150 | -1.491181 |   |
| 19 | 64.089200  | -0.427553 | -1.355529 | -1.160298 | -1.166036 | -1.357600 | -0.138279 |   |

|    | Lag_7     | Lag_8     | Lag_9     | Lag_10    |
|----|-----------|-----------|-----------|-----------|
| 0  | NaN       | NaN       | NaN       | NaN       |
| 1  | NaN       | NaN       | NaN       | NaN       |
| 2  | NaN       | NaN       | NaN       | NaN       |
| 3  | NaN       | NaN       | NaN       | NaN       |
| 4  | NaN       | NaN       | NaN       | NaN       |
| 5  | NaN       | NaN       | NaN       | NaN       |
| 6  | NaN       | NaN       | NaN       | NaN       |
| 7  | -1.313935 | NaN       | NaN       | NaN       |
| 8  | -2.350474 | -1.534397 | NaN       | NaN       |
| 9  | -1.273009 | -2.102512 | -1.378408 | NaN       |
| 10 | 1.561306  | -1.551484 | -1.652281 | -1.574531 |
| 11 | -1.280714 | 0.232160  | -1.369586 | -2.688745 |
| 12 | 0.656299  | -1.541392 | -1.157122 | -1.604177 |
| 13 | -1.277262 | -1.595567 | -1.362201 | 2.580865  |
| 14 | -1.435055 | -1.531047 | -1.092181 | -1.619347 |
| 15 | -1.358402 | -1.784142 | -1.375195 | -1.800530 |
| 16 | -1.274456 | -1.432213 | -1.448149 | -1.572870 |
| 17 | -1.416235 | -1.553786 | -1.340615 | -1.989255 |
| 18 | -1.272676 | -1.612184 | -1.371604 | -1.556037 |
| 19 | -1.278688 | -1.555441 | -1.454254 | -1.601690 |

```
[32]: fig, axs = plt.subplots(10)
      for i in range(10):
          axs[i].scatter(true_values[i],pred_values[i])
          axs[i].xlabel = 'True Values'
          axs[i].ylabel = 'Predicted Values'
          axs[i].set_title("Lag = "+str(i+1))
          axs[i].set_aspect('equal', adjustable='box')
```



```
[33]: # SVM - Sigmoid Kernel
```

```
[34]: true_values = []
      pred_values = []
      rmse_values = []
      for lag in range(1,11):
```

```

# forming features using lag
y = df['ISMR'][lag:]
X = df.iloc[:-1*lag,: ]

# intialising SVM regressor object and fitting on data
svr = SVR(kernel = 'sigmoid', C=1.0, epsilon=0.1)
svr.fit(X,y)

#generating predictions
preds = svr.predict(X)

# calculating error to help as compare lag performance
rmse = np.sqrt(mean_squared_error(y, preds))

# appending values to a list for later use
true_values.append(y)
pred_values.append(preds)
rmse_values.append((rmse))

```

```

[35]: rmse_table = pd.DataFrame(list(np.array(rmse_values).
→reshape((len(rmse_values),-1)).T))
rmse_table.columns = ["Lag_"+str(i)) for i in range(1,11)]

```

```

[36]: rmse_table

```

```

[36]:      Lag_1      Lag_2      Lag_3      Lag_4      Lag_5      Lag_6  \
0  48.170946  43.842853  47.899698  44.426647  44.279512  45.048021

      Lag_7      Lag_8      Lag_9      Lag_10
0  47.645228  44.34741  46.269744  44.8154

```

```

[37]: table = pd.DataFrame()
table['True ISMR'] = df['ISMR']

```

```

[38]: for i in range(len(pred_values)):
      pred_lag = list(pred_values[i])
      for i in range(df.shape[0]-len(pred_lag)):
          pred_lag.insert(0,np.NaN)
      table['Lag_'+str(i+1)] = pred_lag

```

```

[39]: table[0:20]

```

```










[39]:      True ISMR      Lag_1      Lag_2      Lag_3      Lag_4      Lag_5  \
0   10.243800         NaN         NaN         NaN         NaN         NaN
1   44.161100  34.027572         NaN         NaN         NaN         NaN
2    4.768070 -12.875960  28.951401         NaN         NaN         NaN
3  -63.587700  21.048718 -13.407772  32.951564         NaN         NaN

```

|    |            |            |            |            |            |            |
|----|------------|------------|------------|------------|------------|------------|
| 4  | 12.577200  | 60.093893  | 17.356594  | -18.257679 | 28.537653  | NaN        |
| 5  | -41.577200 | 36.731199  | 51.752396  | 20.128548  | -16.188800 | 30.483062  |
| 6  | -9.546880  | 15.293926  | 31.329154  | 66.722448  | 16.788553  | -15.400557 |
| 7  | 25.370500  | -24.471651 | 12.404151  | 35.520899  | 54.457499  | 18.459257  |
| 8  | 18.277400  | 28.329901  | -23.661365 | 18.679901  | 30.951597  | 57.221047  |
| 9  | 2.221670   | 37.118598  | 23.368752  | -25.000500 | 13.421923  | 32.912049  |
| 10 | 16.886600  | 12.498256  | 31.451549  | 25.803351  | -24.815825 | 14.904802  |
| 11 | 0.132196   | 37.820835  | 9.644611   | 35.455205  | 22.333600  | -24.343484 |
| 12 | -10.537500 | 4.794311   | 32.161871  | 11.677827  | 30.951085  | 23.945288  |
| 13 | -50.820200 | -26.119343 | 2.681259   | 36.336673  | 8.950194   | 32.791229  |
| 14 | -8.513200  | 35.753004  | -25.077095 | 3.967644   | 31.647107  | 10.361638  |
| 15 | -28.769000 | -22.582378 | 30.467775  | -26.550503 | 1.816906   | 33.572105  |
| 16 | -5.404050  | -13.548241 | -21.960223 | 40.638309  | -26.250825 | 3.089745   |
| 17 | -17.658400 | -14.740350 | -13.288244 | -23.060209 | 32.257567  | -25.817090 |
| 18 | 14.971800  | -29.133364 | -14.911261 | -12.087316 | -23.043549 | 34.247210  |
| 19 | 64.089200  | 38.017523  | -27.549879 | -15.320347 | -13.570283 | -22.518325 |

|    | Lag_6      | Lag_7      | Lag_8      | Lag_9      | Lag_10     |
|----|------------|------------|------------|------------|------------|
| 0  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 1  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 2  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 3  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 4  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 5  | NaN        | NaN        | NaN        | NaN        | NaN        |
| 6  | -30.973000 | NaN        | NaN        | NaN        | NaN        |
| 7  | 14.127581  | -30.986977 | NaN        | NaN        | NaN        |
| 8  | -19.264606 | 19.276019  | -33.014871 | NaN        | NaN        |
| 9  | -58.567938 | -18.919056 | 11.075932  | -29.849935 | NaN        |
| 10 | -33.347328 | -64.835179 | -20.351968 | 18.494421  | -29.128892 |
| 11 | -16.180935 | -33.374085 | -55.418856 | -17.281655 | 16.914528  |
| 12 | 22.290031  | -18.956920 | -35.695479 | -59.137275 | -17.467653 |
| 13 | -24.808506 | 24.208640  | -13.280236 | -32.439370 | -58.340743 |
| 14 | -33.327549 | -23.321561 | 24.145148  | -14.089381 | -31.519519 |
| 15 | -11.431189 | -32.946625 | -27.813914 | 27.184566  | -15.376160 |
| 16 | -34.029476 | -10.755476 | -36.094708 | -23.207012 | 24.008558  |
| 17 | -4.341066  | -33.773578 | -11.939305 | -32.452635 | -22.631951 |
| 18 | 23.729844  | -3.396961  | -36.807630 | -8.903960  | -31.396902 |
| 19 | -35.479544 | 25.660374  | -4.513106  | -33.163886 | -9.636338  |

```
[40]: fig, axs = plt.subplots(10)
      for i in range(10):
          axs[i].scatter(true_values[i], pred_values[i])
          axs[i].xlabel = 'True Values'
          axs[i].ylabel = 'Predicted Values'
          axs[i].set_title("Lag = "+str(i+1))
          axs[i].set_aspect('equal', adjustable='box')
```



```
[41]: #Performing Grid Search for best kernel
```

```
[42]: parameters = {'kernel':('linear', 'poly', 'rbf', 'sigmoid'),'C':  
    ↪(1,2,5,10),'epsilon':(0.1,.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0,1.5,2.0)}  
svr = SVR()  
reg = GridSearchCV(svr, parameters)
```

```
[43]: lag = 1  
  
    # forming features using lag  
y = df['ISMR'][lag:]  
X = df.iloc[:-1*lag,:]  
  
    # performing grid search  
reg.fit(X,y)  
  
    #generating predictions  
preds = reg.predict(X)  
  
    # calculating error to help as compare lag performance  
rmse = np.sqrt(mean_squared_error(y, preds))
```

```
[44]: rmse
```

```
[44]: 31.519244924329566
```

```
[45]: print(reg.best_params_)
```

```
{'C': 2, 'epsilon': 2.0, 'kernel': 'linear'}
```

```
[46]: # We use lag 1 from previous analysis. The best kernel comes out to be linear.
```

```
[ ]:
```